

Subash Konar

Lead Data Engineer

subashkonar13@gmail.com | +918249939823 | Bengaluru, India | subashkonar13.github.io/techblog

SUMMARY

- 10+ years of experience in building scalable data platforms across Azure and GCP
- Expertise in real-time streaming, event-driven architectures, and cloud-native systems
- Hands-on exposure to LLM integration, RAG pipelines, and AI-driven data tooling
- Strong focus on data quality, governance, and AI/ML readiness
- Proven leadership experience managing teams and delivering large-scale projects

PROFESSIONAL EXPERIENCE

Lead Data Engineer

Advisor360 (Finance) | 10/2024 – Present | Bengaluru

- Reduced manual efforts by 60% with Azure Document Intelligence, streamlining tabular data extraction from PDFs.
- Designed a scalable ingestion framework to process multi-schema files, improving data ingestion efficiency.
- Integrated a lightweight RAG (Retrieval-Augmented Generation) pipeline using Azure OpenAI and Azure AI Search to surface contextual answers from ingested financial documents, supporting internal analyst queries — gained practical exposure to chunking strategies, embedding models, and prompt construction.
- Developed standardised knowledge bases for recurring issues, reducing troubleshooting time and enhancing productivity.
- Built robust data pipelines using Orkes and Netflix Conductor, enabling concurrent execution of multiple workflows.
- Segmented data assets into business-enabling domains using TOGAF principles, driving adoption and scalability.
- Enforced data quality requirements, improving data accuracy and establishing monitoring frameworks for timeliness and completeness.
- Deployed batch and event-driven pipelines, facilitating real-time decision-making and reducing processing delays.
- Led cross-functional projects with 100% on-time delivery, collaborating with product teams to define and launch innovative solutions.

Senior Data Engineer

Microsoft (Supply Chain) | 02/2020 – 10/2024 | Bengaluru

- Engineered real-time event-driven data pipelines using PySpark and Azure Event Hub, processing 10M+ events/day into Delta Lake, improving processing efficiency by 40%.
- Designed scalable data platform architectures using Medallion principles (Bronze, Silver, Gold) for batch and streaming workloads.
- Developed an AI-driven data quality solution using LangChain and OpenAI, automating PySpark code generation and reducing validation effort by 60% — this involved working with prompt templates, LLM chains, and output parsers, giving hands-on experience with LLM orchestration patterns.
- Contributed to an internal POC that explored using a RAG-based assistant to help engineers query pipeline documentation and runbooks via natural language; worked on document indexing and retrieval logic using FAISS and Azure Cognitive Search — practical but exploratory in nature.
- Orchestrated 15+ CI/CD pipelines using Azure DevOps and Terraform for infrastructure provisioning and deployment automation.
- Provisioned and managed cloud infrastructure including Azure Synapse, ADLS, and Key Vault using Infrastructure-as-Code.
- Implemented observability using Azure Log Analytics and Grafana, reducing incident response time by 25%.
- Designed and implemented data models (SCD Type 2, CDC pipelines) supporting large-scale analytical workloads.
- Established data governance and lineage using Apache Atlas, ensuring compliance and traceability.
- Led data cleansing and transformation initiatives, reducing redundant data by 97%.
- Mentored 8+ engineers and led Agile/Scrum planning, improving team productivity and delivery consistency.

Datawarehouse Designer

Woolworths (Retail) | 01/2015 – 01/2020 | NSW, Australia

- Migrated enterprise ETL/ELT pipelines from on-prem RDBMS to GCP using Apache NiFi and Airflow, reducing operational costs.
- Built large-scale data processing platform using Scala, Python, Apache Beam, Dataproc, and BigQuery.
- Led migration of 800+ databases from SQL Server to BigQuery, ensuring data consistency and scalability.
- Optimised distributed data processing using PySpark, improving performance by 80%.
- Designed and implemented ML-ready datasets and pipelines for model training using Google AutoML.
- Managed team of 5 engineers, improving productivity by 30% through process optimisation.

Hadoop Developer

Tata Consultancy Services Ltd (Banking) | 11/2012 – 01/2015 | Chennai, India

- Migrated transactional data to Hadoop ecosystem, reducing processing time by 30% and storage costs by 25%.

- Built automated ingestion and archival workflows using Apache NiFi, reducing manual intervention by 90%.
- Ingested data from Oracle and MySQL using optimised pipelines, reducing latency by 15%.
- Developed automation frameworks using shell scripting to improve data pipeline efficiency.
- Built UI components and backend services using Java, improving client usability and reporting capabilities.

SKILLS

Languages: Python, SQL, Scala

Big Data & Orchestration: Apache Spark (Batch & Streaming), Apache Airflow, Apache Kafka, Apache Beam, Apache NiFi, Netflix Conductor, dbt

Cloud Platforms: Microsoft Azure (Synapse, ADLS, Event Hub, Key Vault), GCP (BigQuery, Dataflow, Dataproc, GCS)

Data Warehouses: BigQuery, Snowflake, Azure Synapse Analytics

DevOps & Infra: Terraform, Docker, Kubernetes, Azure DevOps, Helm

Data Modelling: SCD Type 2, Data Vault, Dimensional Modelling (Star/Snowflake), Delta Lake

Architecture: Event-driven (CQRS, Idempotency), CDC, Medallion, Incremental Processing, Schema Evolution, Batch & Real-time Pipelines

Data Governance: Data Quality Frameworks, Apache Atlas, GDPR, PII Handling, Data Observability (Grafana, Log Analytics)

AI / LLM: LangChain, OpenAI / Azure OpenAI, RAG pipelines (FAISS, Azure AI Search), Prompt Engineering (working knowledge), LLM orchestration, Feature Engineering, ML Dataset Preparation, Vertex AI, Amazon SageMaker

Integrations: REST APIs, SaaS Integrations, Segment, Oracle, MySQL, SQL Server, Kafka, Event Hub

BI & Testing: Power BI, A/B Testing

EDUCATION

Bachelor of Technology (Computer Science & Engineering)

IGIT, Sarang (BPUT), Odisha, India | 06/2008 – 06/2012

CERTIFICATIONS

- Cloudera CCA 175
- Fabric Analytics Engineer (DP600)
- Azure Data Engineer Associate (DP-203)
- GCP Data Engineering (Coursera)
- Azure Data Scientist Associate (DP-100)

PROJECTS

Fraud Detection in Financial Transactions

- Built machine learning model using Random Forest achieving 85% accuracy.
- Developed end-to-end pipeline for data preprocessing, feature engineering, and model evaluation.

LINKS & LANGUAGES

Website: <https://subashkonar13.github.io/techblog/> | **LinkedIn:** <https://www.linkedin.com/in/subashkonar/> | **Languages:** English, Spanish, Hindi